
CS 769 Project Presentation

Comparison of Multi-Armed Bandits Policies

Presented by -

Prishat Bachhar - 213050078
Subir Kumar Parida - 21v980019
Yashwant Raghuwanshi - 213050002

Introduction to Online Convex Optimization

K: convex decision set, and **F**: set of convex functions.

A repeated game of T rounds between a player and an adversary, at each round $\mathbf{t} \in \mathbf{T}$, then the OCO framework,

- player chooses a point $\mathbf{x}_t \in \mathbf{K}$
- adversary independently chooses a loss function $\mathbf{f}_t \in \mathbf{F}$
- player suffers a loss $\mathbf{f}_t(\mathbf{x}_t)$ and receives a feedback \mathbf{F}_t

The regret of the setting can be defined as:

$$\text{Regret}_T = \sum_{t=1}^T f_t(x_t) - \min_{x^* \in K} \sum_{t=1}^T f_t(x^*)$$

Objective: To minimize the expected regret

Online Gradient Descent

Algorithm 1: Online Gradient Descent

Input: Convex set $K, T, x_1 \in K$, step sizes $\{\eta_t\}$

```
1 for  $t = 1$  to  $T$  do
2   |   Play  $x_t$  and observe loss  $f_t(x_t)$ 
3   |   Update and Project:
4   |        $y_{t+1} = x_t - \eta_t \nabla f_t(x_t)$ 
5   |        $x_{t+1} = \Pi_K(y_{t+1})$ 
6 end
```

The Exploration-Exploitation Dilemma

- **Exploration** - improve knowledge for long-term benefit
- **Exploitation** - exploit knowledge for short-term benefit

Multi Armed Bandits (MAB)

Algorithm 20 Simple MAB algorithm

```
1: Input: OCO algorithm  $\mathcal{A}$ , parameter  $\delta$ .
2: for  $t = 1$  to  $T$  do
3:   Let  $b_t$  be a Bernoulli random variable that equals 1 with probability
      $\delta$ .
4:   if  $b_t = 1$  then
5:     Choose  $i_t \in \{1, 2, \dots, n\}$  uniformly at random and play  $i_t$ .
6:
7:     Let
        
$$\hat{\ell}_t(i) = \begin{cases} \frac{n}{\delta} \cdot \ell_t(i_t), & i = i_t \\ 0, & \text{otherwise} \end{cases}.$$

8:     Let  $\hat{f}_t(\mathbf{x}) = \hat{\ell}_t^\top \mathbf{x}$  and update  $\mathbf{x}_{t+1} = \mathcal{A}(\hat{f}_1, \dots, \hat{f}_t)$ .
9:   else
10:    Choose  $i_t \sim \mathbf{x}_t$  and play  $i_t$ .
11:    Update  $\hat{f}_t = 0, \hat{\ell}_t = \mathbf{0}, \mathbf{x}_{t+1} = \mathbf{x}_t$ .
12:   end if
13: end for
```


- In this setting the player only knows the loss of the action chosen and not the entire loss gradient.

Other Strategies for MAB

- Epsilon Decay Algorithm
 - Keep reducing the value of epsilon $\frac{1}{(1 + \frac{t}{\log(\text{arms})})^2}$
- UCB Algorithm
 - Reduce the uncertainty associated with an action
 - Done by sampling more number of times
- Thompson Sampling
 - Sample the next action using a Beta distribution of the number of times a reward was achieved and the number of times a loss was achieved by choosing that action.

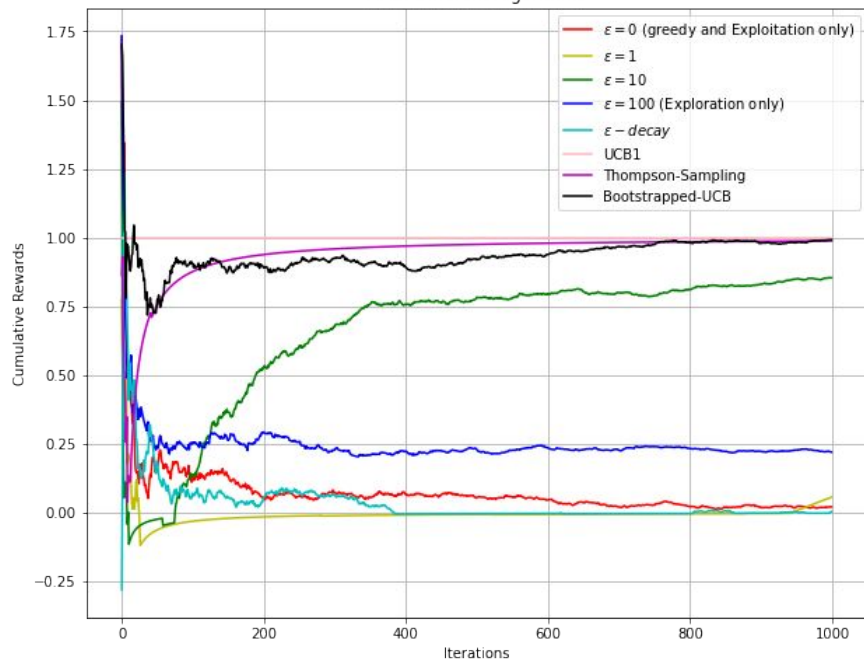
Results

- Higher exploration rate leads to higher average loss over time
- This is due to the fact that the algorithm explores even when it is almost sure of the best action
- The other strategies and their convergence is given in the following slides.




Results

MAB: Different Algorithms



Convergence rates of different algorithms


MAB Algorithms - Arms Selection Strategy



Histogram of actions chosen by the algorithms

Contextual MAB

- It is an extension of multi-armed bandits.
- The algorithm observes a context, makes a decision, choosing one action from a number of alternative actions, and observes an outcome of that decision.



Bibtex Dataset:

- It is a Multilabel text classification for automated tag suggestion:
 - containing tags that people have assigned to different papers
 - the goal being to learn to suggest tags based on features from the papers
- [dataset link](#)

	Obs.	Feats.	Labels
BibTeX	7,395	1,836	159

Epsilon-Greedy for Contextual Bandits:

Inputs probability $p \in (0, 1]$, decay rate $d \in (0, 1]$, oracles $\hat{f}_{1:k}$

- 1: **for** each successive round t with context \mathbf{x}^t **do**
- 2: With probability $(1 - p)$:
- 3: Select action $a = \operatorname{argmax}_k \hat{f}_k(\mathbf{x}^t)$
- 4: Otherwise:
- 5: Select action a uniformly at random from 1 to k
- 6: Update $p := p \times d$
- 7: Obtain reward r_a^t , Add observation $\{\mathbf{x}^t, r_a^t\}$ to the history for arm a
- 8: Update oracle \hat{f}_a with its new history

Implementation:

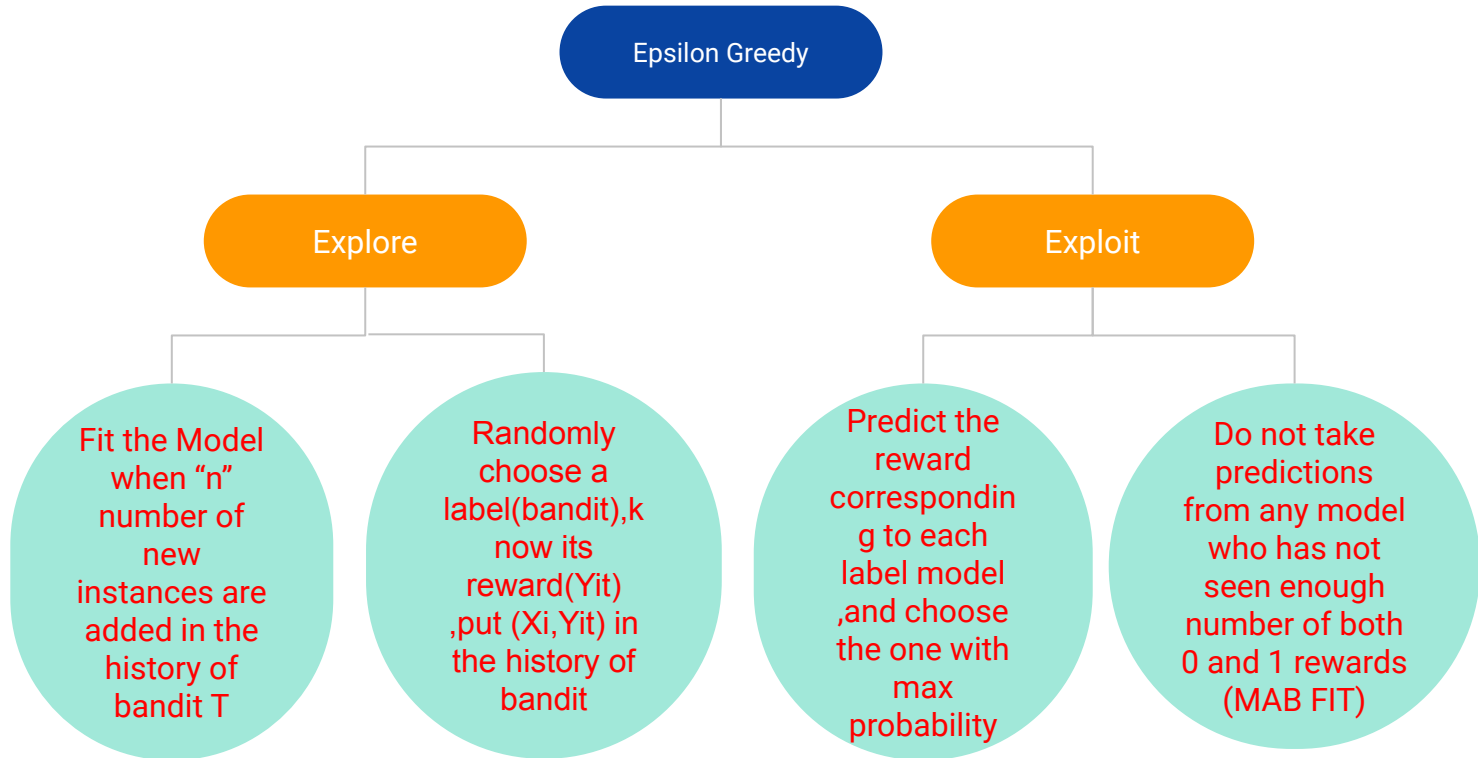
- Firstly, the bibtex dataset is transformed into binary classification data for each label corresponding to each training instance.
- **MAB-FIRST:**

Inputs const. a, b , threshold m , contextual bandit policy π_k , covariates \mathbf{x}

Output score for arm \hat{r}_k

- 1: **if** $|\{r \in \mathcal{R}_k \mid r = 0\}| < m$ or $|\{r \in \mathcal{R}_k \mid r = 1\}| < m$ **then**
 - 2: Sample $\hat{r}_k \sim \text{Beta}(a + |\{r \in \mathcal{R}_k \mid r = 1\}|, b + |\{r \in \mathcal{R}_k \mid r = 0\}|)$
 - 3: **else**
 - 4: Set $\hat{r}_k = \pi_k(x)$
- return** \hat{r}_k
-


WorkFlow:



Parameters being Hypertuned:

- a and b values for beta distribution ($a=3, b=7$)
- Subset of bandits (labels) on which models are being tuned due to sparsity of data (opt subset length = 10)
- Refit the corresponding bandit depending on the number of new training instances being added to its history, a tradeoff against computational power utilized (opt. value = 5)

Results



Critical Analysis

- In the context of MAB, a higher probability of exploration leads to more regret as the algorithm is forced to explore bad choices.
- **ϵ -greedy family:** The algorithms are not at all concerned towards obtaining the true mean of an arm for exploitation
- **UCB, Thompson:** The algorithms tries to reduce the uncertainty about the true reward of an arm
- **UCB:**
 - UCB eliminates the randomness of the ϵ -greedy family algorithms
 - The algorithm gives the benefit of doubt to a less explored arm at a certain time t , and explores it to identify the true distribution/reward
- **Thompson:**
 - Tries to select the arm which has previously given more number of high rewards.
 - It tries to understand the true distribution of the rewards associated with the arms.

Thank you